



# Artificial Intelligence Techniques (MSc)

*Lecture 2: Adversarial Search and Game Playing*  
Delft University of Technology

Dr. Nils Bulling

September 20, 2016

# Outline

- 1 Games and Search: Introduction
- 2 Optimal Decision Making
- 3 Non-Optimal Decision Making: Resource Constraints
- 4 Stochastic Games
- 5 Partially Observable Games
- 6 Summary

**Lecturer:** Dr. Nils Bulling

**Teaching Assistant:** Chris Rozemuller, MSc

**Course email:** ai@ii.tudelft.nl

**Lecture:**

- Wednesday, September 21st, 08:45-10:45

**Tutorial** and **Test:**

- Tutorial: Thursday, September 22nd
- Test: Thursday, September 29th

Please note the following points

- 1 The **exercises** and **assignments** are very **important** for the understanding and for passing the exam!
- 2 I would appreciate it to hear about any **flaws**, **typos**, **comments** and any other feedback which helps to improve the lecture.

# Reading Material I

- The slides are quite detailed and should in general be sufficient.
- General reading :



Russel, S. and Norvig, P. (2010).  
*Artificial Intelligence: a Modern Approach.*

Prentice Hall, 3 edition. Chapter 4.5, 5

# Acknowledgement and Copyright

The lecture is based on and uses material of:



Russel, S. and Norvig, P. (2010).

*Artificial Intelligence: a Modern Approach.*

Prentice Hall, 3 edition. Chapter 4.4 and 5.

Pictures and material are taken from this source. The **copyright** of the material taken stays with the **authors and publisher of the original material. It is prohibited to copy or to distribute these slides in any form.**

The author gratefully acknowledges material and slides provided by **Carholijn Jonker** developed in previous editions of this course at TU Delft.

# Next Section

- 1 Games and Search: Introduction
- 2 Optimal Decision Making
- 3 Non-Optimal Decision Making: Resource Constraints
- 4 Stochastic Games
- 5 Partially Observable Games
- 6 Summary

## Topics for today: **algorithms for playing (board) games**

- formalization of **games as search problems**
- finding **optimal decisions** in games (e.g Tic-Tac-Toe)
- complexity: **pruning** and **heuristic methods** (e.g Chess)
- **partially observable games** (e.g Poker)
- **stochastic games** (e.g Backgammon)



# Games

- **setting:**
  - a player's outcome depends on actions of others
  - often **self-interested players**
- **unpredictability** of other players and **incomplete information**  
⇒ **contingencies** How to deal with them?
- Players use **strategies**. A strategy is a **conditional plan** that specifies how to reply to an opponent.
- Can we use classical search algorithms like **breath-first** and **depth-first search** to find optimal strategies?
- Extension of classical search algorithms: **adversarial search** (and-or search) to find winning strategies.
- Problem: Huge **states-spaces** often prevent exact solutions  
⇒ **clever techniques** and **heuristics methods** are needed

# Brief Historical Overview

1846 (Babbage): discussion on computer's solving games

1912/44 (Zermelo/van Neumann/Morgenstern):

algorithm for perfect play (Minimax), foundation of game theory

1945-1950 (Zuse/Wiener/Shannon):

chess: finite horizon, approximate evaluation

1948 (Turing) chess program (paper machine)

1952/59 (Samuel) checkers: machine learning to improve heuristic

1955 (McCarthy) pruning of search tree to allow deeper search

1997 (IBM) IBM's Deep Blue wins over chess world champion Garry Kasparov

# Properties of Games

## Example 1.1 (Chess)

Which “game properties” has chess?


- How many **players**?
- Which **information** do players have about the **state of the game**?
- Is the outcome of an action **completely deterministic**?

We can classify games along the following dimensions:

<b>Types of Games</b>	deterministic	stochastic
perfect information	chess, checkers, go, othello	backgammon, monopoly
imperfect information	battleship, blind tictactoe, Kriegspiel	bridge, poker, scrabble, nuclear war

In the following, we often consider **2-player** (**Max** and **Min**),  
**turn-based**, **deterministic**, **strictly competitive**, **perfect information**

The (finite) rules of a game induce a **game tree**, consisting of all possible plays (is it always finite?):



We give a **semi-formal definition** of an **adversarial search problem**.

### Definition 1.2 (Formulation as Search Problem)

An **adversarial search problem** consists of the following elements:


- $s_0$ : **initial (game) state** of the game
- $\text{Player}(s)$ : defines the **player** whose **turn** it is at state  $s$
- $\text{Action}(s)$ : returns set of **legal actions** at state  $s$
- $\text{Result}(s, a)$ : (**game**) **state** which **results** if action  $a$  is being executed in  $s$
- $\text{TerminalTest}(s)$ : returns **true** if the **game is over**; false otherwise.
- $\text{Utility}(s, p)$ : returns the **utility** of player  $p$  at **terminal node**  $s$ .

The **game tree** of the search problem is the **tree induced** by functions **Action** and **Result** in the initial state  $s_0$ . We often use  $S$  to denote the set of all game states.

Note the differences between: **adversarial search problem**, **game tree** and a **search tree** over the adversarial search problem!

# Notation of ply

The execution of each action in a game is also called a **ply**.  
For example, the following game tree consists of **two plies**, or **2-pplies**.



In chess a **move** consists of **2-pplies**.

# Next Section

- 1 Games and Search: Introduction
- 2 Optimal Decision Making**
- 3 Non-Optimal Decision Making: Resource Constraints
- 4 Stochastic Games
- 5 Partially Observable Games
- 6 Summary

# Computing Optimal Decisions

How to find an **optimal strategy**?

How to **search** through the **game tree**? Can we use standard search algorithms like **depth-first search** or **breadth-first search**?

We discuss the following:

- **optimal decision making**: find the optimal strategy of a player  $\rightsquigarrow$  **Minimax algorithm**
- improve performance of Minimax algorithm  $\rightsquigarrow$  **alpha-beta Minimax**



# Subsection I

## 2 Optimal Decision Making

### Minimax Algorithm


Minimax Algorithm in Multi-Player Games

Alpha-Beta Pruning

Consider the following 2-player game tree. Players are denoted by **Max** and **Min**. **Terminal nodes** are the leaf nodes; they are labelled with Max's utility value. What is the **optimal** strategy of Max?

MAX

MIN



⇒ **FeedbackFruits** Idea: We assume that the **opponent plays optimally** and compute the best possible action under this assumption.

$$\text{Minimax}(s) = \begin{cases} \text{Utility}(s) & \text{if } \text{TerminalTest}(s) \\ \max_{a \in \text{Actions}}(\text{Minimax}(\text{Result}(s, a))) & \text{if } \text{Player}(s) = \text{Max} \\ \min_{a \in \text{Actions}}(\text{Minimax}(\text{Result}(s, a))) & \text{if } \text{Player}(s) = \text{Min} \end{cases}$$

$\text{Minimax}(s)$  recursively computes the utility of the optimal action at state  $s$ . It is called the **minimax value** at  $s$ .

### Exercise 2.1 (Non-optimal opponents)

*Minimax assumes that the opponent plays optimally. Suppose the opponent does not play optimally. Show the following.*

- 1 For non-optimally playing opponents the Minimax value is at least as good as if played against an optimal opponent.
- 2 For non-optimally playing opponents, the Minimax value may not be optimal (i.e. give a concrete game tree as counterexample).

The **Minimax algorithm** performs a **depth-first search** through the game tree. We assume that Max is the starting player.

**function** MINIMAX-DECISION(*state*) *returns an action*  
    **return**  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$

---


**function** MAX-VALUE(*state*) *returns a utility value*  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow -\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
    **return** *v*

---

**function** MIN-VALUE(*state*) *returns a utility value*  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow \infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
         $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
    **return** *v*

## Example 2.1 (Minimax)

Apply the Minimax algorithm to the following game tree. Compute the minimax value and the optimal action:



# Properties of Minimax

## Definition 2.2 (Branching factor)

The **branching factor** at a node is the number of legal moves at that node. The **maximal branching factor** is the maximal branching factor at some node. The **average branching factor** is the average branching factor across all nodes.

We will often use **branching factor** to refer to the average branching factor.

## Example 2.3

What is the average branching factor of the following games?

- TicTacToe  $\rightsquigarrow 4$
- Backgammon  $\rightsquigarrow \sim 400$ , maximal one can be much higher 400
- Chess  $\rightsquigarrow 35$
- Go  $\rightsquigarrow 250$

Recall: A **search strategy** prescribes how a state space/game tree is searched.

### Definition 2.4 (Completeness, optimality)

A **search strategy** is called

- **complete**, if it **finds a solution**, provided there exists one at all.
- **optimal**, if whenever it produces an output, this output is an **optimal solution**, i.e. one with the best utility.

It is assumed that you know the following classical **search strategies and their properties**:

- depth-first search
- breath-first search
- iterative deepening search
- etc.

We analyse the **time complexity** of the Minimax algorithm. How many nodes does the Minimax algorithm examine?

Suppose the game tree has the (finite) **maximal branching factor**  $b$  and a **depth** of at most  $m$  **levels** (starting from 0), i.e. an  $m$ -ply game tree.

Minimax performs a full depth-first search:  $\rightsquigarrow$  FeedbackFruits

**Time complexity** :  $b + b^2 + \dots + b^m = \frac{1-b^{m+1}}{1-b} - 1 = O(b^m)$

**Space complexity** :  $O(bm)$

Can be improved to  $O(m)$  if only one successor is kept in memory, which is then currently updated.

**Complete** : yes (if  $m$  is **finite**)

**Optimal** : yes, against an **optimal opponent**



# Exact values Do Not Matter

The exact value of the terminal nodes of a game tree do not matter as long as it is ensured, that the relative **ordering is kept**

↪ **ordinal utility function**.

Let  $t$  be a **transformation function** on terminal nodes to real numbers such that for all **terminal nodes**  $n_1$  and  $n_2$  with **values**  $x, y \in \mathbb{R}$ , respectively, in the original game tree it holds that:

- $x \leq y$  iff  $t(n_1) \leq t(n_2)$ , and
- $x < y$  iff  $t(n_1) < t(n_2)$ .

Then, the actions computed by the Minimax algorithm on the original game tree and to the game tree obtained by updating the utility values according to the transformation function  $t$  are identical.

## Exercise 2.2 (Exact values do not matter)

*State the above more formally and give a formal proof.*

# Subsection I

## 2 Optimal Decision Making


Minimax Algorithm

Minimax Algorithm in Multi-Player Games

Alpha-Beta Pruning

We have considered a **two player setting**. This can be generalized to **arbitrary players** and **non-zero sum games**:

to move  
A



### Exercise 2.3 (Minimax in multi-player turn-based games)

*Extend the Minimax algorithm to the multi-player setting and compute the optimal strategy of player A in the game tree shown above.*

# Subsection I

## 2 Optimal Decision Making

Minimax Algorithm

Minimax Algorithm in Multi-Player Games

Alpha-Beta Pruning

The Minimax algorithm computes an **optimal strategy**, but it has to **search the whole tree**. This is often not possible.

### Example 2.5 (Chess)

How many nodes to search in Chess up to the following depth:


- 1-ply:  $b = 35$
- 2-ply (1 move):  $b^2 = 1225$
- 3-ply:  $b^3 = 42875$
- 4-ply (2 moves):  $b^4 = 1.5$  millions

The number of nodes to search grows too fast—**exponential in the depth**. Can we **prune** the search tree and at the same time **retain optimality**?

Yes, using the following idea:

- **Prune subtrees which cannot yield optimal solutions.**
- In order to decide which subtrees to prune, store the value of the **best “choices”** ( $\alpha$  resp.  $\beta$ ) of the players (Max resp. Min) encountered so far.

The soundness of the pruning follows from the observation that the minimax value is independent of the minimax value of pruned subtrees:



$$\begin{aligned} \text{Minimax}(s_0) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \\ &= \max(3, z, 2) \quad \text{for } z = \min(2, x, y) \leq 2 \\ &= 3 \end{aligned}$$

The **Alpha-Beta Minimax algorithm** implements the idea of pruning, by updating **two values for each node**:

**$\alpha$ -value**: value of **best action** for *Max* found so far at any choice point **along the path** (i.e. highest value)

**$\beta$ -value**: value of **best choice** for *Min* found so far at any choice point **along the path** (i.e. lowest value)

In the search tree, the values are represented as interval  $[\alpha, \beta]$ .

Now, pruning works as follows:

- 1 **Update the values** along a path.
- 2 If the computed value of the **current node** would give a **worse solution for the other player** then **prune remaining subtrees**. (As the other player would not choose this node.)

Suppose it is **Min's turn** and it can ensure a value  $v$  with  $v \leq \alpha$ . Then, Max has a **different action at least as good**  $\rightsquigarrow$  **prune** the tree.

**function** ALPHA-BETA-SEARCH( $state$ ) **returns** an action  
     $v \leftarrow \text{MAX-VALUE}(state, -\infty, +\infty)$   
    **return** the *action* in ACTIONS( $state$ ) with value  $v$







---

**function** MAX-VALUE( $state, \alpha, \beta$ ) **returns** a *utility value*  
    **if** TERMINAL-TEST( $state$ ) **then return** UTILITY( $state$ )  
     $v \leftarrow -\infty$   
    **for each**  $a$  **in** ACTIONS( $state$ ) **do**  
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
        **if**  $v \geq \beta$  **then return**  $v$   
         $\alpha \leftarrow \text{MAX}(\alpha, v)$   
    **return**  $v$




**function** MIN-VALUE( $state, \alpha, \beta$ ) *returns a utility value*  
**if** TERMINAL-TEST( $state$ ) **then return** UTILITY( $state$ )  
 $v \leftarrow +\infty$   
**for each**  $a$  **in** ACTIONS( $state$ ) **do**  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    **if**  $v \leq \alpha$  **then return**  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
**return**  $v$

## Example 2.7 (Alpha-Beta Pruning)



## Proposition 2.1 (Correctness of Alpha-Beta Minimax)

Let  $P \in \{\text{Max}, \text{Min}\}$  and  $O$  be the opponent. Let  $n$  be a node in the game tree and suppose that  $P$  has an *action to reach  $n$*  in the next step. If  $P$  has a *better action  $m$  along the path* to  $n$  (i.e. either at the parent node of  $n$  or at any choice point further up in the path (preventing from reaching  $n$ )), then *node  $n$  will never be reached* if  $P$  plays a strategy induced by the Minimax algorithm.



Proof.

↪ Exercise




# Performance

The improve in the **performance** is sensitive to the **action ordering** (i.e. the order in which the game tree is traversed).

## Example 2.8 (Order important for efficiency)

In Example 2.7, (e) and (f) no terminal node can be pruned. What if the 3rd successor of  $D$  had been the first child of  $D$ ?



## Remark 2.1 (Alpha-Beta Minimax)

- “optimal action ordering”: time complexity  $\approx O(b^{m/2}) = O(\sqrt{b^m}) \rightsquigarrow$  *Exercise*
- Thus, Alpha-Beta Minimax can *examine a game tree roughly twice as deep* as Minimax.
- “random generation of action”: time complexity  $\approx O(b^{3m/4})$
- Conclusion: a *good ordering is very important* in order to search large game trees.

## Exercise 2.4 (Alpha-Beta Minimax with optimal move ordering)

Show that in the case of an “optimal action ordering” the Alpha-Beta Minimax algorithm has a worst-case time complexity of  $O(b^{m/2})$  where  $m$  is the maximal depth of the game tree.

# Next Section

- 1 Games and Search: Introduction
- 2 Optimal Decision Making
- 3 Non-Optimal Decision Making: Resource Constraints**
- 4 Stochastic Games
- 5 Partially Observable Games
- 6 Summary

So, far we considered **optimal decisions**. Often, there are **resource limits** that do not allow to reach optimal decisions.

### Example 3.1 (Search under time constraints)

Suppose we play a **chess** and decisions must be made within **100 seconds**. Moreover, the computer can **explore  $10^4$  nodes** per second. (Recall, the average branching factor of chess is  $b \approx 35$ .)

- How many plies can *MiniMax* (optimally) explore? About 4:  
 $O(b^m) \leq 10^6 \rightsquigarrow 35^{3.9} \approx 10^6 \rightsquigarrow m \approx 4$   
(for the exact value, solve:  $\frac{1-b^{m+1}}{1-b} \leq 10^6$ )
- ... and alpha-beta search (with a good action ordering)? About 8
- This is considered the level of a **good chess player**.

What if **time limit** but **no terminal node** is reached? How to make a decision? Possibly a **non-optimal one**?

# Heuristic Minimax Value

Instead of a **cutoff** after a predefined number of steps, we consider a quality-based approach.

- A **cutoff test** replaces the **terminal test** in the Minimax algorithm. It decides when to stop the search, possibly using an **evaluation function**.
- An **evaluation function** determines how good the current state is. It returns a **heuristic utility value**.

The **heuristic Minimax value** is defined as follows:

$$H\text{-Minimax}(s) = \begin{cases} Eval(s) & \text{if } CutoffTest(s, d) \\ \max_{a \in Actions} (H\text{-Minimax}(Result(s, a))) & \text{if } Player(s) = Max \\ \min_{a \in Actions} (H\text{-Minimax}(Result(s, a))) & \text{if } Player(s) = Min \end{cases}$$

The **CutoffTest**( $s$ ) may also take into consideration the current depth  $d$ .



## ③ Non-Optimal Decision Making: Resource Constraints

### Evaluation Functions

Minimax with Cutoffs

Deterministic Games in Practice

An **evaluation function** estimates the worth of the current state. This is also how **human chess players** evaluate the current state of the game. How to define such a function ?

**Basic requirements** on the evaluation function:

- 1  $Eval(s)$  should give the same value as  $Utility(s)$  for **terminal states**  $s$  (e.g. don't turn win into loose).
- 2 Evaluation of the function must **not be too complex**: time limits!
- 3 At **non-terminal states**, the evaluation value should be **strongly correlated** with the **chance of winning**.

Note that **uncertainty** is introduced by **computational limitations** not by the game itself.

How to define **good evaluation functions**? This is a difficult task. For example, player experience, static features, sub-solutions.

## Basic construction of an evaluation function:

- 1 Define **features**  $f_1, \dots, f_n$  of a state which induce a vector-based evaluation function  $f(s) = (f_1(s), \dots, f_n(s))$ .
- 2 Values of features induce **equivalence classes**, called **categories**, of states. Same values of features corresponds to same category, e.g.

$$C_1 = \{s \mid f_1(s) = c_1^1, \dots, f_n(s) = c_n^1\}, \quad \text{here } c_j^1 \text{ are fixed}$$

- 3 States in a category are usually of different quality. Define the evaluation function for a category as the **expected value** of the **states in the category**, e.g.

$$\mathbf{E}(C_1) = 0.7 \cdot \mathbf{1} + 0.1 \cdot \mathbf{0} + 0.2 \cdot (-\mathbf{1})$$

(where **feature value** 1, 0, -1 corresponds to **win**, **draw**, **loose**, respectively).

- 4 Define the **evaluation of a state** as the evaluation of its category:

$$\text{Eval}(s) = \mathbf{E}(C) \quad \text{if } s \in C$$

### Remark 3.1 (Evaluation function)

- Often, Minimax allows to consider 4-ply lookahead if classes of categories/features is not too big, and if the features can easily be computed.
- The construction requires a lot of experience and many categories; often too many to be useful.

In practise: Compute **numerical values from each feature** and **combine them** into a single value (independent of the notation of category).

### Example 3.2 (Evaluation of a Chess state)

In chess, e.g., a queen could be assigned a worth of 10 and a rook 6.

Another feature could be defined as:

$$f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$$



(a) White to move



(b) White to move

- One could use a **linear weighted sum** of features ( $w_i$  are weights):  

$$y = \sum_{i=1}^n w_i x_i$$

$$Eval(s) = w_1 f_1(s) + \dots + w_n f_n(s) = \sum_{i=1}^n w_i f_i(s)$$

This assumes that features are **independent** of each other. Therefore, **non-linear** functions can be used as well.

# Subsection I

## ③ Non-Optimal Decision Making: Resource Constraints

Evaluation Functions

Minimax with Cutoffs

Deterministic Games in Practice

Is it rather straight-forward how to incorporate evaluation functions into **Minimax** and **Alpha-Beta Minimax**. The *TerminalTest* in the Alpha-Beta Minimax algorithm is simply replaced by:


if *CutoffTest*( $s, d$ ) then return *Eval*( $s$ )

where  $d$  represents the **current depth**.

A depth limit is used to ensure that the **answer is given within a given time limit**. How to determine a good depth limit?

- Simple approach: introduce a **fixed depth limit  $d$** .
- More robust: combine with **iterative deeping depth-first search**  $\rightsquigarrow$  when time is up, return value of deepest **completed search**

This simple implementation can lead to errors. **Crucial changes** in the player situation need to be **taken into account**. Let us reconsider the following situation:



White to move

- Seems to be a **winning situation** for **Black**.
- But: White can capture the black queen, indicating a **winning situation** for White.
- It is said that the state of the game is **not quiescent**.



Some notes:

- A situation is **quiescent** if an action cannot yield fundamental changes in the state evaluation.  
     $\rightsquigarrow$  **Cutoffs should only be made in quiescent states.**
- More improvements are possible/necessary: e.g. horizon effect/singular extensions

# Subsection I

## ③ Non-Optimal Decision Making: Resource Constraints

Evaluation Functions

Minimax with Cutoffs

Deterministic Games in Practice

# Overview: Deterministic Games in Practice

**Checkers:** Jonathan Schaeffer's Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a **precomputed endgame database** defining **perfect play for all positions involving 8 or fewer pieces** on the board, a total of 444 billion positions. 2007: **solved**:

[http://www.newscientist.com/article/dn12296-checkers-solved-after-years-of-number-c.html#.VHzc8\\_mG98E](http://www.newscientist.com/article/dn12296-checkers-solved-after-years-of-number-c.html#.VHzc8_mG98E)

Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997,

<http://www.youtube.com/watch?v=3EQA679DFRg>.

Deep Blue searches 200 million positions per second and 12-ply, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 plys.

**Null move heuristic:** shallow search where opponent can move twice at the beginning (gives a lower bound)


Go: Difficult for Computers. Problem: **branching factor**  $> 300$ , so most programs use pattern knowledge bases to suggest plausible actions.

Recent progress: In 2016 AlphaGo (Google) won against a professional Go master (9 dan)

# Next Section

- 1 Games and Search: Introduction
- 2 Optimal Decision Making
- 3 Non-Optimal Decision Making: Resource Constraints
- 4 Stochastic Games**
- 5 Partially Observable Games
- 6 Summary

So far, we considered **deterministic** and fully observable games.  
What if a game includes **dice throws** or **random moves**?



# Subsection I

## 4 Stochastic Games

### Game Trees with Chance Moves


### Minimax Algorithm for Stochastic Games

We extend game trees with **chance moves**. This can be seen as another player who selects actions **randomly**.

MAX

CHANCE

MIN



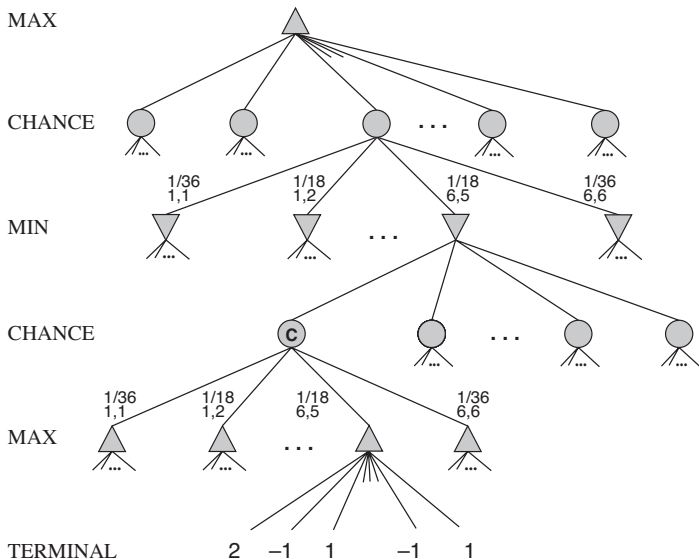
Sources of randomness:

- dice throws
- card shuffling
- coin flipping
- drawing pieces
- etc.

Attention: We **merge the actions** of the **player** and the associated **chance move** into a **single ply** (here Min's action). This affects the branching factor: in the above example, it is  $b = 4$  (rather than  $b = 2$ ) for **Min's ply**.



Backgammon has the following excerpt of a game tree. The situation is that Max moved its pieces, thereafter, Min throws the dices:



# Subsection I

## 4 Stochastic Games

Game Trees with Chance Moves

Minimax Algorithm for Stochastic Games

# Expected Minimax Value


$$\text{Expectiminimax}(s) = \begin{cases} \text{Utility}(s) & \text{if } \text{TerminalTest}(s) \\ \max_{a \in \text{Actions}} (\text{Expectiminimax}(\text{Result}(s, a))) & \text{if } \text{Player}(s) = \text{Max} \\ \min_{a \in \text{Actions}} (\text{Expectiminimax}(\text{Result}(s, a))) & \text{if } \text{Player}(s) = \text{Min} \\ \sum_{r \in \sigma} (\text{Pr}(r) \text{Expectiminimax}(\text{Result}(s, r))) & \text{if } \text{Player}(s) = \text{Chance} \end{cases}$$

where

- $\sigma$  represents a **stochastic action** (like a dice role) and  $r \in \sigma$  a possible **outcome** of the action
- $\text{Result}(s, r)$  is the same state as  $s$ , together with the **information that the outcome of the stochastic action is  $r$** .

The **ExpectiMiniMax algorithm** is defined analogously.

As before, **evaluation functions** can be used to define a cut-off version of the Expectiminimax algorithm. But note, that with chance moves the evaluation is **sensitive to the actual values**:



Here, the optimal action is **different** in both game trees.  
 Behavior preserved if **positive linear transformations** are considered.

## Remark 4.1 (Practical Aspects)

- Chance nodes increase significantly the *branching factor*  $b$ 
  - **time complexity**:  $O(b^m n^m)$  where  $n$  is number of **distinct dice rolls**  $\rightsquigarrow$  *FeedbackFruits*
  - **Backgammon**: 21 is the number of *distinct dice rolls* (with two dice);  $\approx 20$  moves, *average branching factor* about 420 (but much more with doubles, e.g. 1-1 means: move 4 pieces each 1 positions  $\rightsquigarrow$  branching factor of 4000 and more)
  - *4-plies*:  $(21 \cdot 20)^4 \approx 3.1 \cdot 10^{10}$  nodes
  - limited value in look-ahead for increased depth
  - *alpha-beta pruning* less effective (**Why?**)
- TD-Gammon Gammon (G. Tesauro, 1992, IBM)
  - *artificial neural net* trained by a form of temporal-difference learning
  - *Depth 2 search + very good evaluation function*: world-champion level

## Exercise 4.1 (Alpha Beta Pruning for Stochastic games)

*Give an intuitive argument why alpha-beta search may not be very effective in stochastic games.*

# Next Section

- 1 Games and Search: Introduction
- 2 Optimal Decision Making
- 3 Non-Optimal Decision Making: Resource Constraints
- 4 Stochastic Games
- 5 Partially Observable Games**
- 6 Summary

In games like chess and backgammon, players have **perfect information** about the world. There are many games in which players have **incomplete information**; they cannot see the complete state of the game. Examples are:

- poker
- bridge
- scrabble
- etc.

How can we model these games appropriately?

We consider **belief states** rather than actual states of the game. Such states **encodes all situation an agent cannot distinguish from the current situation**.



A **belief state**  $b$  consists of all **physical states** the player considers possible, together with a probability.

### Definition 5.1 (Belief state)

Let  $S$  be a set of **game states** of a game. A **belief states**  $b$  over  $S$  is a **probability distribution** over  $S$ .

We also represent a **belief state**  $b$  as a subset  $X = \{s_1, \dots, s_n\} \subseteq S$  where  $b$  is defined as follows (clearly not all belief states can be represented in such a way):

$$b(s) = \begin{cases} \frac{1}{n} & \text{if } s \in X \\ 0 & \text{otherwise.} \end{cases}$$

In that case we also identify  $b$  with  $X$ . How is the **(expected) utility of a belief state** defined?


$$\text{Utility}(b) = \sum_{s \in S} b(s) \cdot \text{Utility}(s)$$

## Example 5.2 (Vacuum world)


The **vacuum world** is defined as follows:

- It consists of **two rooms**.
- Each room may be **dirty**.
- A vacuum agent can **move** from one room to the other, or **suck the dirt**. Actions are *Left*, *Right*, *Suck*.
- The vacuum agent **does not know its location**, nor can it see whether a **room is dirty or not**.


The **goal** is to ensure that **both rooms are clean**. A game state:




The game states of the vacuum world are:




The transitions in the physical world look at follows:



What is a **belief state** in the world? Each of the two game states in a belief state can be associate **probability 0.5**. A state transition in the belief space (which action is being performed?):



The following picture shows the reachable part of the belief space.



- In partially observable games, **strategies** are defined over belief states, or rather **sequences of belief states**/ **sequences of percepts**.
- A **winning strategy** must yield a belief state in which all physical states are winning.
- These games can also be combined with **non-deterministic actions**, **chance moves**, etc.
- In general the **reachable belief states** are **exponential** wrt. the set of game states. Finding optimal solutions becomes even more challenging.

### Example 5.3 (Vacuum world)

Is there a winning strategy (**sure winning**) in the vacuum world, i.e. one in which **both rooms are clean** and the **robot knows this**?  $\rightsquigarrow$

FeedbackFruits

Action sequence: Left, Suck, Right, Suck

## Example 5.4 (Kriegspiel)

**Kriegspiel** is a partially observable variant of chess.

- Players only **see their own pieces**.
- They announce their actions and a referee announces whether each move is **legal** or **illegal**.
- The **referee also informs about captures, checks, checkmates or stalemates**, among other things.

So, the player **learns from** the **announcements** of the referee.


A **winning strategy** corresponds to a strategy which ensures a **guaranteed checkmate**, one which leads to a checkmate for every possible board state in the **current belief state**.



black only has  
a **king left**

figure encode  
**belief states**

strategy to  
find out  
about **true  
position of  
king**



Partially observable games offers **different notions of winning strategies**. For example, in Kriegspiel:

- **guaranteed checkmate**: guaranteed winning
- **probabilistic checkmate**: winning **with probability 1**. E.g. finding the lonely king on the board. (In general this is not always possible.)
- **accidental checkmate**: player wins but could not **know** this from its belief state.
- It may **not always be optimal to perform the “best action”** as it may provide too much information  $\rightsquigarrow$  concept of **equilibrium**

As another example of partially observable games, we discuss **card games**. Cards are **dealt randomly** at the beginning of the game, and are of course **invisible to other players**. How could we compute the best action? **Now how can we solve such games?**

What about the following idea: does it work?

- 1 Consider all **possible deals**  $s_1, \dots, s_r$  of the invisible cards.
- 2 Compute the **optimal action**  $a_i$  wrt. to each  $s_i$  assuming that  $s_i$  describes the actual distribution of the cards. Use (a variant of) Minimax.
- 3 **Take the action** that **maximizes the expected utility** wrt. the games induced by  $s_i$ : (note that  $a$  is the same action for all  $s_i$ !)

$$\operatorname{argmax}_a \sum_{i=1}^r \Pr(s_i) \operatorname{Minimax}(\operatorname{Result}(s_i, a))$$

What can be problematic with this approach?

Note, that we assume that the game is **fully observable** for each **fixed guessed deal  $s$** . Thus, we average over the result of the game if we knew the correct state. Therefore, it is also called **averaging over clairvoyance**.

The following example illustrates why **averaging over clairvoyance does not work** for (many) games with incomplete information.

## Example 5.5 (Problem with “averaging over clairvoyance”)

- Day 1:
- **Road A** leads to a heap of gold; **Road B** leads to a fork
  - **left fork**: bigger heap of gold; **right fork**: you will be run over by a bus.
- Day 2:
- **Road A** leads to a heap of gold; **Road B** leads to a fork
  - **right fork**: bigger heap of gold; **left fork**: you will be run over by a bus.
- Day 3:
- **Road A** leads to a heap of gold; **Road B** leads to a fork
  - **one fork**: bigger heap of gold; **other fork**: you will be run over by a bus.
  - You **cannot distinguish the two forks**.

What happens if **averaging over clairvoyance** is applied?  $\rightsquigarrow$

FeedbackFruits

- In averaging over clairvoyance, ***B* is taken** because:
  - *B* is the best choice at Day 1.
  - *B* is the best choice at Day 2.
  - *B* is the best choice as Day 3 as it is either Day 1 or Day 2.

The possible “deals” are **instantiated**, and Road *B* is always the best move, followed by **different actions at the fork**.

- In day three the **belief state** of the agent should have been considered and not an (**instantiated**) **game state**. Note that the flawed approach assigns **different actions** in states which the agent **cannot distinguish**.
- The **averaging over clairvoyance** approach always yields states of **perfect knowledge** which is not always appropriate.
- The right way to solve partially observable games is to consider **belief states**  $\rightsquigarrow$  later in the course

# Next Section

- 1 Games and Search: Introduction
- 2 Optimal Decision Making
- 3 Non-Optimal Decision Making: Resource Constraints
- 4 Stochastic Games
- 5 Partially Observable Games
- 6 Summary**

# Summary and Conclusions

- Solving games is challenging
- and important, e.g. it plays a role in verification tasks (multi-agent systems)
- Solving games is often (theoretically) possible, but in practice, limited resources prevent optimal solutions.
- Many “simple” games have a huge branching factor which makes them hard to solve (e.g. Go); so far...
- Heuristic approaches are needed.
- Partially observable and stochastic games are even harder to solve.
- Some games with incomplete information are impossible to solve algorithmically: they are undecidable.



# References I



Russel, S. and Norvig, P. (2010).  
*Artificial Intelligence: a Modern Approach*.  
Prentice Hall, 3 edition.

# References I



Russel, S. and Norvig, P. (2010).  
*Artificial Intelligence: a Modern Approach*.  
Prentice Hall, 3 edition.